Rachel Chasin

https://www.linkedin.com/in/rachel-chasin/
rachel.chasin@gmail.com | 201-289-3525

EDUCATION

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

M.Eng. in Electrical Engineering and Computer Science
June 2013 | Cambridge, MA
GPA: 5.0 / 5.0

S.B. in Mathematics and Computer Science and Engineering
June 2012 | Cambridge, MA
GPA: 5.0 / 5.0

TECHNOLOGIES

Languages and Frameworks
Python • uWSGI • Pyramid • Numpy/Scipy/Sklearn/Pandas
Java • Jetty • Jackson • Dropwizard
JavaScript (familiar) • jQuery • AngularJS

Datastores
MySQL • Cassandra • Redshift / PostgreSQL

Distributed Systems
Spark • EMR • Kafka • S3

DevOps
AWS CLI/APIs/console, various services • Puppet (familiar) • Terraform (familiar)

Development Environment
Unix • git

SELECTED EXPERIENCE

YELP | Search Team | Software Engineer
August 2017 – present | San Francisco, CA

- Created MVP then fully productionized system to predict business properties from keyword mentions in reviews.
- Used learning-to-rank to train regression models for ranking businesses, released and A/B tested the models online.
- Migrated unspecialized searches from deprecated homegrown system to elasticsearch.
- Improved team tools: ranking model training pipeline (better data persistence, optional use of Spark); qualitative recall/ranking comparison tool.

YELP | Spam and Abuse Detection Team | Software Engineer
August 2013 – August 2017 | San Francisco, CA

- Developed and maintained Yelp's centralized service that synchronously detects disruptive or abusive content submitted by users, using a DSL of rules.
  - Major contributor to initial implementation, deployment, migration of existing functionality, validation, documentation, logging, and monitoring.
  - Advised on subsequent major changes to the service, including an asynchronous mode and experimentation support.
- Developed and productionized specialized review quality classifier for businesses without much content.
- Improved infrastructure for spammy or low quality review detection.
  - Introduced Redshift to store data too large for MySQL, cluster soon became a data warehouse and new computation engine for the team.
  - Led a project using Kafka to asynchronously move values of signals as they are computed into multiple datastores (Redshift, Cassandra).
  - Reduced redundant computation of signals when no new data is available.

GOOGLE | Risk Engineering Team | Software Engineering Intern
June 2012 - August 2012 | Mountain View, CA

- Working at the boundary between a payment fraud detection service and its clients, created a minimal in-memory client fallback for use upon service failure
- Published risk engine responses to an internal publish/subscribe system, with multiple subscribers from other team applications

RESEARCH

MIT COMPUTER SCIENCE AND ARTIFICIAL INTELLIGENCE LAB | Clinical Decision Making Group | Research Assistant
September 2011 – June 2013 | Cambridge, MA

- Natural language processing research in word sense disambiguation techniques on medical domain text (clinical notes), culminating in a paper and an M.Eng. thesis.
- Published as: R. Chasin, A. Rumshisky, Ö. Uzuner, P. Szolovits. 2014. Word sense disambiguation in the clinical domain: a comparison of knowledge-rich and knowledge-poor unsupervised methods. *J Am Med Inform Assoc.* http://jamia.bmj.com/content/early/2014/01/17/amiajnl-2013-002133. 2014 Jan 17.